

# Advanced Static Analysis of Atomicity in Concurrent Programs through Facebook Infer

Master's Thesis

Bc. Dominik Harmim


Supervisor: prof. Ing. Tomáš Vojnar, Ph.D.

xharmi00@stud.fit.vutbr.cz

Brno University of Technology, Faculty of Information Technology



14th June 2021


-  Detecting and checking desired **atomicity** of function call sequences.
  - Often required in concurrent programs.
  - Violation** may cause nasty errors.

```
void invoke(char *method) {
    ...
    if (server.is_registered(method)) {
        server.invoke(method);
    }
    ...
}
```

The sequence of **is\_registered** and **invoke** should be **executed atomically**.

If **not locked**, the method can be unregistered by a concurrent thread.

- Open-source **static analysis framework** for **interprocedural analyses**.
  - Based on **abstract interpretation**.
- Highly **scalable**.
  - Follows principles of **compositionality**.
  - Computes function **summaries** **bottom-up** on call-trees.
- Supports C, C++, Java, Obj-C, C#.

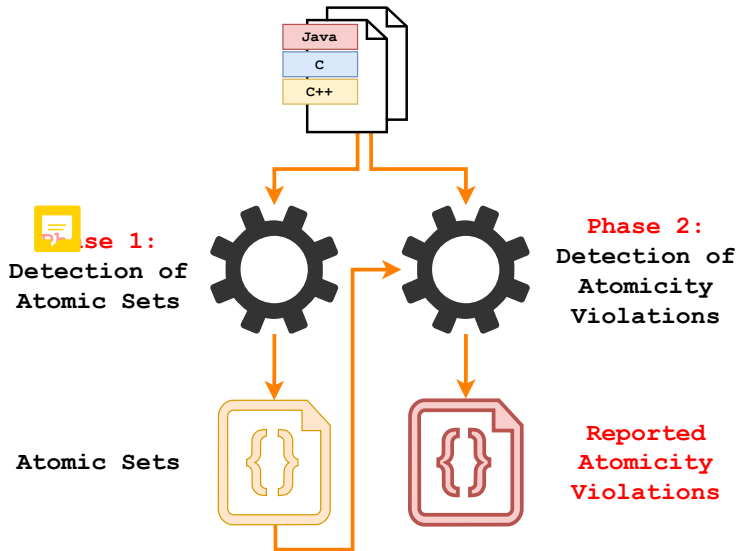


- Facebook Infer plugin created within the author's BSc thesis:



HARMIM, D. *Static Analysis Using Facebook Infer to Find Atomicity Violations*. Brno, 2019. Bachelor's thesis. Brno University of Technology, Faculty of Information Technology. Supervisor VOJNAR, T.

- **Assumption:** call sequences executed atomically once should (probably) be executed always atomically.
- Implemented for C programs that use PThread locks.
- Limited scalability on extensive codebases.
- Reports many false alarms when analysing real-life code.





## 1 Detection of **atomic call sets**.

- Approximates sequences by sets.
- Summary:**  $\chi \in 2^\Sigma \times 2^{2^\Sigma}$   
(set of all calls, set of atomic call sets)

```
void f() {
    lock(L);
    x(); y(); z(); // x.y.z -> {x,y,z}
    unlock(L);
    a();
    lock(L);
    z(); y(); x(); // z.y.x -> {x,y,z}
    unlock(L);
}
```

$$\chi_f = (\{a, x, y, z\}, \{\{x, y, z\}\})$$

$$\chi'_f = (x \cdot y \cdot z \cdot a, \{x \cdot y \cdot z, z \cdot y \cdot x\})$$

## 2 Detection of **atomicity violations**.

- Derives "**atomic pairs**" from the first phase:  $\Omega \in 2^{\Sigma \times \Sigma}$
- Looks for **non-atomic pairs** of calls assumed to **run atomically**.
- Summary:**  $\chi \in 2^{\Sigma \times \Sigma}$   
(set of atomicity violations)

```
void g() {
    a(); x(); y(); b();
}
```

$$\Omega = \{(x, y), (x, z), (y, x), (y, z), (z, x), (z, y)\}$$

$$\Omega' = \{(x, y), (y, z), (z, y), (y, x)\}$$

$$(x, y) \in \Omega \implies \chi_g = \{(x, y)\}$$

- Support for C++ and Java.
  - Working with advanced locks: re-entrant locks, monitors, lock guards, etc.
- Distinguishing different lock instances.
  - Approximating lock objects using syntactic access paths—a representation of heap locations via the paths used to access them.
- Analysis's parametrisation:
  - ignoring generic functions versus concentrating on critical functions;
  - limiting the number of calls or the depth of nested calls in critical sections.

- **Scalability** evaluated on 54 real-life complex C programs.
  - 806,431 LOC in total.
- ~~Double acceleration~~ in average.

	v1.0.0		v2.0.0	
	Phs. 1	Phs. 2	Phs. 1	Phs. 2
<b>Avg. Time (s)</b>	70.98	109.11	37.96	50.93
<b>Total Time (s)</b>	4,117	5,892	2,164	2,750

- Experiments with **Apache Cassandra** and **Apache Tomcat** (both ~250 KLOC).
  - Successfully **rediscovered** already fixed reported real bugs.
  - The number of reported bugs was **significantly reduced** ( $\sim 4\times$ ).
  - Still hard to say which of the bugs are real — the **accuracy** needs to be further improved.



- Proposed and implemented extensions for Atomer:
  - approximation with sets, support for C++ and Java, distinguishing different lock instances, parametrisation of the analysis.
- Successfully tested and experimentally evaluated.
  - Both scalability and accuracy were significantly increased.
- Experiments with real-life programs.

## Future goals

- Further increase accuracy/reduce the number of false alarms.
  - Combining with dynamic analysis.
  - Statistic ranking of atomic functions/reported errors.
  - Considering formal parameters of functions.
  - Machine learning of analysis' parameter values.

---

<sup>1</sup>The preliminary results of this work were presented at the Excel@FIT'21 (it won two awards). It is supported by the H2020 ECSEL project VALU3S.



1 Plánujete podniknout další kroky pro zařazení **Atomeru** do **hlavní větve** frameworku **Facebook Infer**?

- **Ano**, určitě bychom se rádi o zařazení v budoucnu pokusili.
- Repositář Atomeru je **pravidelně aktualizován na nejnovější verzi** frameworku.
- Atomer už byl dříve (úspěšně) **prezentován** a **konzultován** s **vývojáři Inferu**.

