

Advanced Static Analysis of Atomicity in Concurrent Programs through Facebook Infer

SEP — Term Project

Dominik Harmim

Supervisor: prof. Ing. Tomáš Vojnar, Ph.D.

xharmi00@stud.fit.vutbr.cz

Brno University of Technology, Faculty of Information Technology



3rd February 2021


- Detecting and checking desired **atomicity of call sequences**:
 - Often required in **concurrent programs**.
 - **Violation** may cause **nasty errors**.

```
void invoke(char *method) {  
    ...  
    if (server.is_registered(method)) {  
        server.invoke(method);  
    }  
    ...  
}
```

The **sequence** of **is_registered** and **invoke** should be **executed atomically**.

If **not locked**, the method can be unregistered by a **concurrent thread**.

- An open-source **static analysis framework** for **interprocedural analyses**.
 - Based on **abstract interpretation**.
- Highly **scalable**:
 - Follows principles of **compositionality**.
 - Computes function **summaries** **bottom-up** on call trees.
- Supports C, C++, Java, Obj-C (and C#).

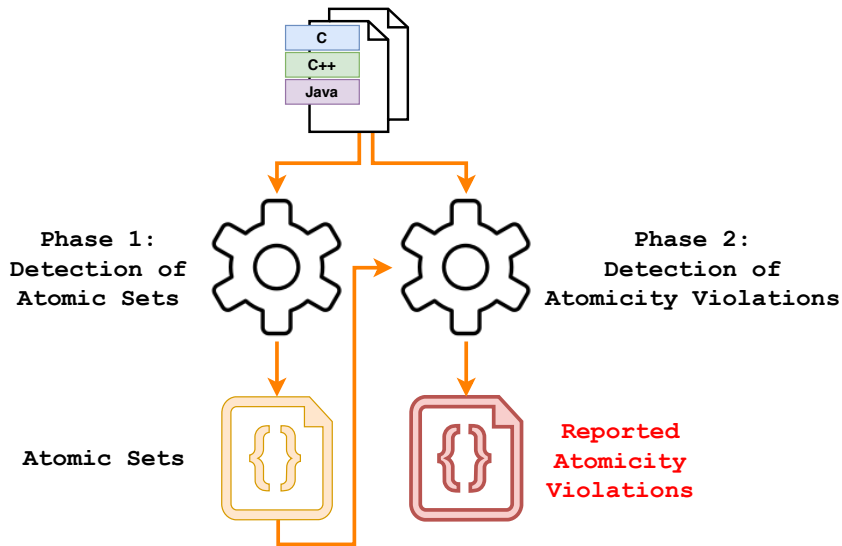


- A **Facebook Infer plugin** created within the bachelor's thesis:



HARMIM, D. *Static Analysis Using Facebook Infer to Find Atomicity Violations*. Brno, 2019. Bachelor's thesis. Brno University of Technology, Faculty of Information Technology. Supervisor VOJNAR, T.

- **Assumption:** **Call sequences** executed **atomically once** should (probably) be executed **always atomically**.
- Implemented for **C** programs that use **PThread locks**.
- Limited **scalability** on extensive codebases.
- Reports many **false alarms** when analysing **real-life** code.



1 Detection of atomic call sets.

- Approximates sequences by sets.
- Summaries:** (set of all calls, set of atomic call sets)

```
void f() {
    lock(L);
    x(); y(); z();
    unlock(L);
    a();
    lock(L);
    z(); y(); x();
    unlock(L);
}
```

$Summary_f: (\{a, x, y, z\}, \{\{x, y, z\}\})$

$Summary'_f: ((x, y, z, a), \{(x, y, z), (z, y, x)\})$

2 Detection of atomicity violations.

- Looks for non-atomic pairs of calls assumed to run atomically.
- Summaries:** (set of first calls, set of last calls, set of atomicity violations)

```
void g() {
    a();
    x(); y();
    b();
}
```

$AtomPairs_g: \{(x, y), (x, z), (y, x), (y, z), (z, x), (z, y)\}$

$AtomPairs'_g: \{(x, y), (y, z), (z, y), (y, x)\}$

$Summary_g: (\{a\}, \{b\}, \{(x, y)\})$

- Support for **C++ and Java**:
 - **C++**: `std::mutex`, `std::lock`, `std::lock_guard`, ...
 - **Java**: `monitors` (`synchronized`), `Lock`, `ReentrantLock`, ...
- Distinguishing **multiple (nested) locks** used:
 - Approximates lock objects using **syntactic access paths**—a representation of heap locations via the paths used to access them.
- **Parametrisation** of the analysis by a user:
 - ignore `generic functions`/concentrate on `critical functions`,
 - limit the number of calls or the depth of analysing nested calls in **critical sections**.

- The **correctness** of the extensions was verified on **hand-crafted** programs.
- **Real-life Java** programs — **Apache Cassandra** and **Tomcat** (~200k LOC).
 - Successfully **rediscovered** already fixed reported real bugs.
 - So far quite some **false alarms** — need to further increase **accuracy**.

Extensions for Atomer:

- Proposed and implemented extensions for Atomer:
 - approximation of sequences by sets, support for C++ and Java, distinguishing multiple locks, parametrisation of the analysis.
- Successfully tested and experimentally evaluated.
- Experiments with real-life programs.

Future goals:

- Further analysis of real-life programs with an effort to find and report new bugs.
- Increase accuracy/reduce the number of false alarms:
 - Support for interprocedural locks.
 - Combination with dynamic analysis.
 - Statistic ranking of atomic functions/reported errors.