

Advanced Static Analysis of Atomicity in Concurrent Programs through Facebook Infer

Master's Thesis

Dominik Harmim

Supervisor: prof. Ing. Tomáš Vojnar, Ph.D.

xharmi00@stud.fit.vutbr.cz

Brno University of Technology, Faculty of Information Technology



24th June 2021

The goal: improve detection of atomicity violations within Facebook Infer.

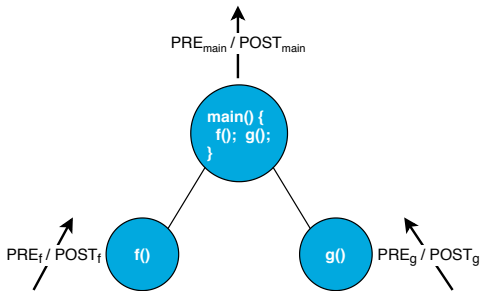
- Detecting and checking desired atomicity of function call sequences.
 - Often required in concurrent programs.
 - Violation may cause nasty errors.

```
void invoke(char *method) {  
    ...  
    if (server.is_registered(method)) {  
        server.invoke(method);  
    }  
    ...  
}
```

The sequence of **is_registered** and **invoke** should be executed atomically.

If not locked, **method** can be unregistered by a concurrent thread.

- Open-source **static analysis framework** for **interprocedural analyses**.
 - Based on **abstract interpretation**.
- Highly **scalable**.
 - Follows principles of **compositionality**.
 - Computes function **summaries** **bottom-up** on call-trees.
- Supports C, C++, Java, Obj-C, C#.



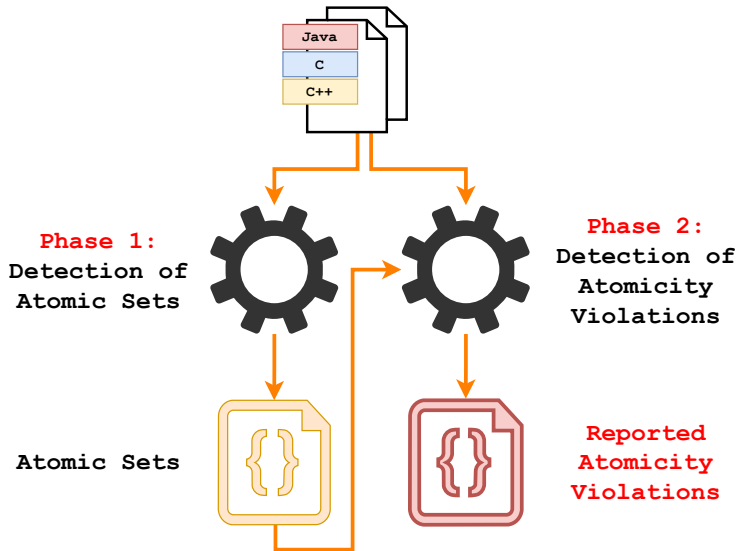
- Facebook Infer plugin created within the author's BSc thesis:



HARMIM, D. *Static Analysis Using Facebook Infer to Find Atomicity Violations*. Brno, 2019. Bachelor's thesis. Brno University of Technology, Faculty of Information Technology. Supervisor VOJNAR, T.

- **Assumption:** call sequences executed atomically once should (probably) be executed always atomically.
- Implemented for C programs that use PThread locks.
- Limited scalability on large codebases.
- Reports many false alarms when analysing real-life code.

- Approximating **sequences** of calls by **sets** of calls (described later on).
- Support for **C++** and **Java**.
 - Working with **advanced locks**: re-entrant locks, monitors, lock guards, etc.
- Distinguishing **different lock instances**.
 - Approximating **lock objects** using **syntactic access paths** — a representation of **heap locations** via the paths used to access them.
- Analysis's **parametrisation**:
 - **ignoring** **generic functions** versus **concentrating on critical functions**;
 - **limiting** the **number of calls** or the **depth of nested calls** in **critical sections**.



1 Detection of atomic call sets.

- Approximates sequences by sets.
- Summary:** $\chi \in 2^{2^\Sigma}$ (set of atomic call sets)

```
void f() {
    lock(L);
    x(); y(); z(); // x.y.z -> {x,y,z}
    unlock(L);
    a();
    lock(L);
    z(); y(); x(); // z.y.x -> {x,y,z}
    unlock(L);
}
```

$$\chi_f = \{\{x, y, z\}\}$$

2 Detection of atomicity violations.

- Derives "atomic pairs" from the first phase: $\Omega \in 2^{\Sigma \times \Sigma}$.
- Looks for non-atomic pairs of calls assumed to run atomically.
- Summary:** $\chi \in 2^{\Sigma \times \Sigma}$ (set of atomicity violations)

```
void g() {
    a(); x(); y(); b();
}
```

$$\Omega = \{(x, y), (x, z), (y, x), (y, z), (z, x), (z, y)\}$$

$$(x, y) \in \Omega \implies \chi_g = \{(x, y)\}$$

1 Detection of atomic call sets.

- Approximates sequences by sets.
- Summary:** $\chi \in 2^{2^\Sigma}$ (set of atomic call sets)

```
void f() {
    lock(L);
    x(); y(); z(); // x.y.z -> {x,y,z}
    unlock(L);
    a();
    lock(L);
    z(); y(); x(); // z.y.x -> {x,y,z}
    unlock(L);
}
```

$$\chi_f = \{\{x, y, z\}\}$$

2 Detection of atomicity violations.

- Derives "atomic pairs" from the first phase: $\Omega \in 2^{\Sigma \times \Sigma}$.
- Looks for non-atomic pairs of calls assumed to run atomically.
- Summary:** $\chi \in 2^{\Sigma \times \Sigma}$ (set of atomicity violations)

```
void g() {
    a(); x(); y(); b();
}
```

$$\Omega = \{(x, y), (x, z), (y, x), (y, z), (z, x), (z, y)\}$$

$$(x, y) \in \Omega \implies \chi_g = \{(x, y)\}$$

- **Scalability** evaluated on 54 real-life complex C programs.

- 806,431 LOC in total.

- On average, **twice faster**.

	v1.0.0		v2.0.0	
	Phs. 1	Phs. 2	Phs. 1	Phs. 2
Avg. Time (s)	70.98	109.11	37.96	50.93
Total Time (s)	4,117	5,892	2,164	2,750

- Experiments with **Apache Cassandra** and **Apache Tomcat** (both ~ 250 KLOC).
 - Successfully **rediscovered** already fixed reported real bugs.
 - The number of reported bugs was **significantly reduced** ($\sim 4\times$).
 - Still hard to say which of the bugs are real — the **accuracy** needs to be further improved.

- Proposed and implemented extensions for Atomer:
 - approximation with sets, support for C++ and Java, distinguishing different lock instances, parametrisation of the analysis.
- Successfully tested and experimentally evaluated.
 - Both scalability and accuracy were significantly increased.
- Experiments with real-life programs.

Future goals


- Further increase accuracy/reduce the number of false alarms.
 - Combining with dynamic analysis.
 - Statistic ranking of atomic functions/reported errors.
 - Considering formal parameters of functions.
 - Machine learning of analysis' parameter values.


¹The preliminary results of this work were presented at the Excel@FIT'21 (won two awards). It is supported by the H2020 ECSEL project VALU3S.

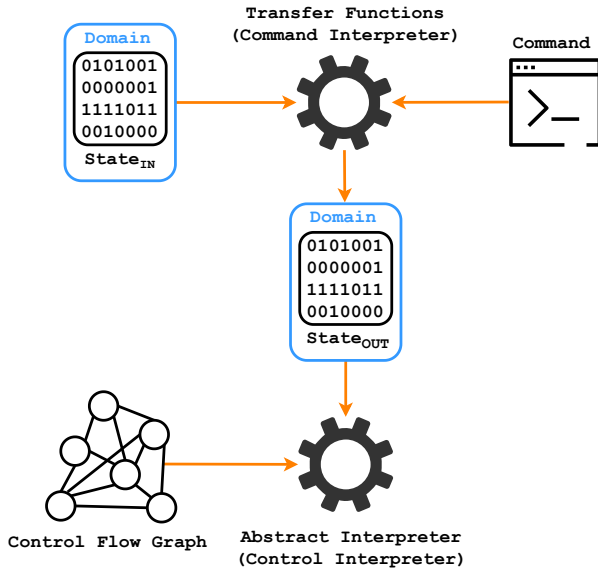
1 Plánujete podniknout další kroky pro zařazení **Atomeru** do **hlavní větve** frameworku **Facebook Infer**?

- **Ano**, určitě bychom se rádi o zařazení v budoucnu pokusili.
- Repositář Atomeru je **pravidelně aktualizován na nejnovější verzi** frameworku.
- Atomer už byl dříve (úspěšně) **presentován** a **konsultován** s **vývojáři Inferu**.
 - Presentace na **Infer Practitioners Workshop** v rámci konference **PLDI 2020**.

- Access path used for a lock's identification: $\pi \in \Pi ::= \text{Var} \times \text{Field}^*$,
 - Var is a set of all variables,
 - Field is a set of field names.
- Identification of a critical section: $(\pi, l) \in \Pi \times \mathbb{N}^\top$,
 - π is an access path that identifies a lock object that locks the section,
 - l is the number of locks of the lock object identified by π ,
 - \mathbb{N}^\top denotes $\mathbb{N} \cup \{\top\}$,
 - \top represents a number larger than some upper bound $t \in \mathbb{N}$.
- Representation of a lock guard: $(\pi_g, L) \in \Pi \times 2^\Pi$,
 - π_g is an access path that identifies the lock guard,
 - L is a set of access paths that identify lock objects associated with the guard.







Real-life bug in a package `org.apache.catalina.core.StandardContext`

```
public void addParameter(String name, String value) {  
    ...  
    if (parameters.get(name) != null)  
        throw new IllegalArgumentException  
            (sm.getString("standardContext.parameter.duplicate", name));  
  
    // Add this parameter to our defined set  
    synchronized (parameters) {  
        parameters.put(name, value);  
    }  
    fireContainerEvent("addParameter", name);  
}
```